Long short-term memory

Neural Networks that remember

Victor ADASCALITEI

Neural Networks recap

Neural Networks recap

What we know so far?

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

Generative Adversarial Networks (**GAN**s)

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

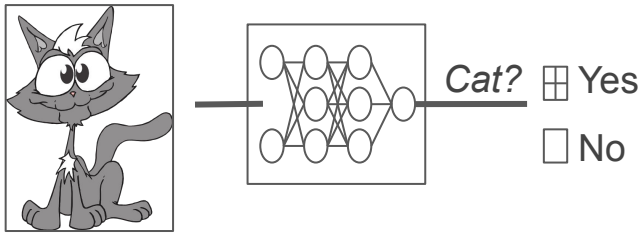

Generative Adversarial Networks (**GAN**s)

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

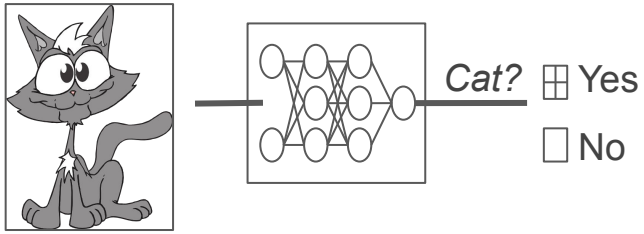

Generative Adversarial Networks (**GAN**s)

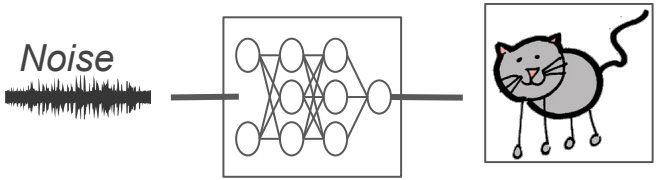

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

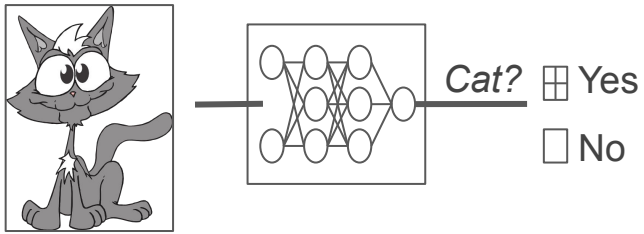

More generally

Generative Adversarial Networks (**GAN**s)

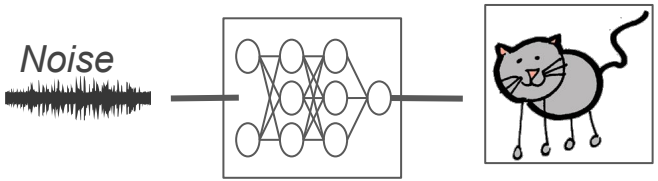

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

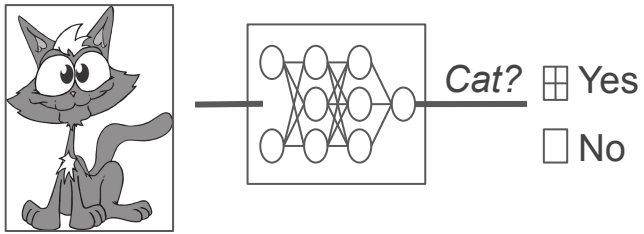

Generative Adversarial Networks (**GAN**s)

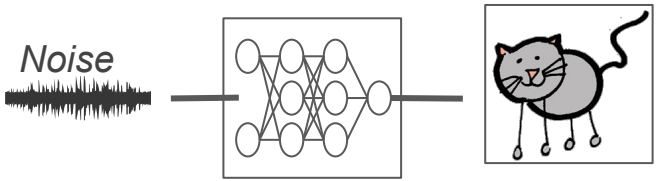

More generally

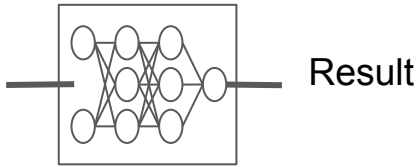

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

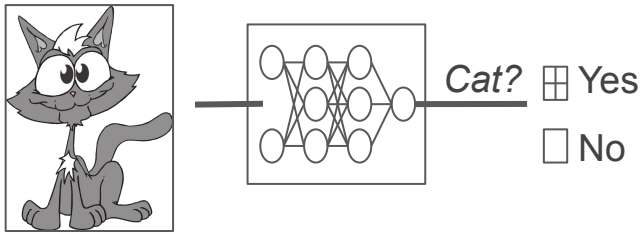

Generative Adversarial Networks (**GAN**s)

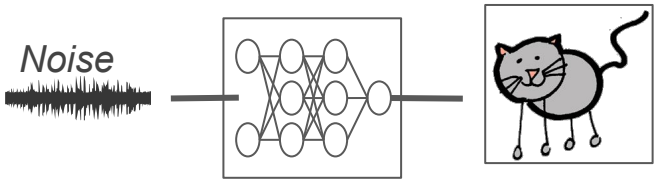

More generally

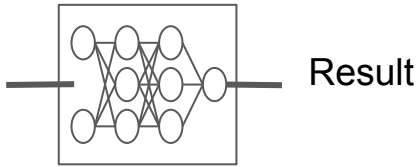

OR

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

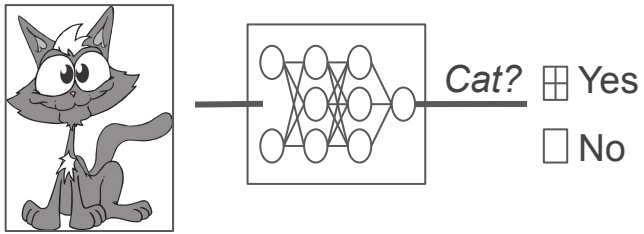

Generative Adversarial Networks (**GAN**s)

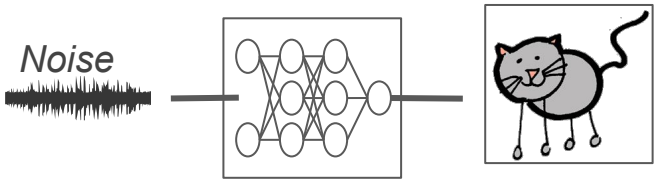

More generally

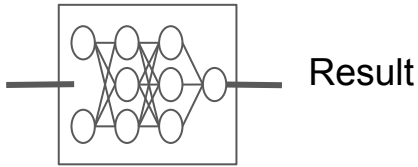

OR

Neural Network(Data) = Result

Neural Networks recap

What we know so far?

Convolutional Neural Networks (**CNN**s)

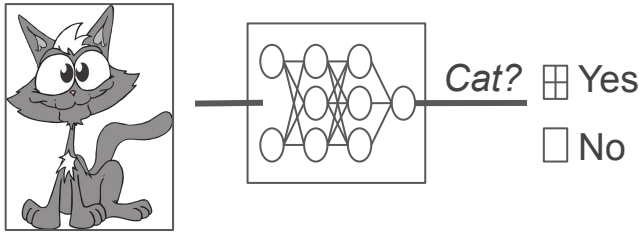

Generative Adversarial Networks (**GAN**s)

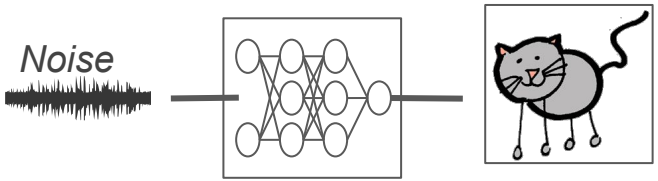

More generally

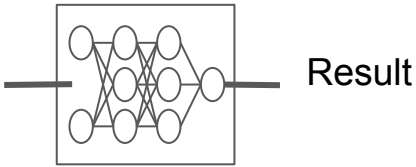

OR

Neural Network(Data) = Result

$f(x) = y$

Neural Networks recap

They are universal function approximators

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images and classifiers

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images and classifiers

but

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images and classifiers

but

What about text ?

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images and classifiers

but

What about text ?

… and speech ?

… and music ?

Neural Networks recap

They are universal function approximators

(Given **enough** data, they can approximate any function)

True for images and classifiers

but

What about text ?

… and speech ?

… and music ?

What about time/context dependent data?

Neural Networks

Input in relation with the output

Neural Networks

Input in relation with the output

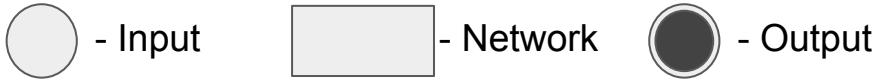

Neural Networks

Input in relation with the output

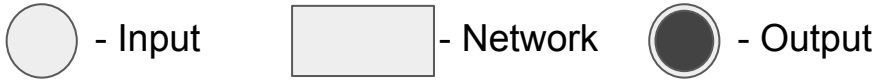

Neural Networks

Input in relation with the output

One to One

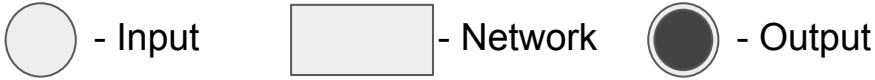

Neural Networks

Input in relation with the output

One to One

One to Many

Many to One

Many to Many

- Input

- Network

- Output

Neural Networks

Input in relation with the output

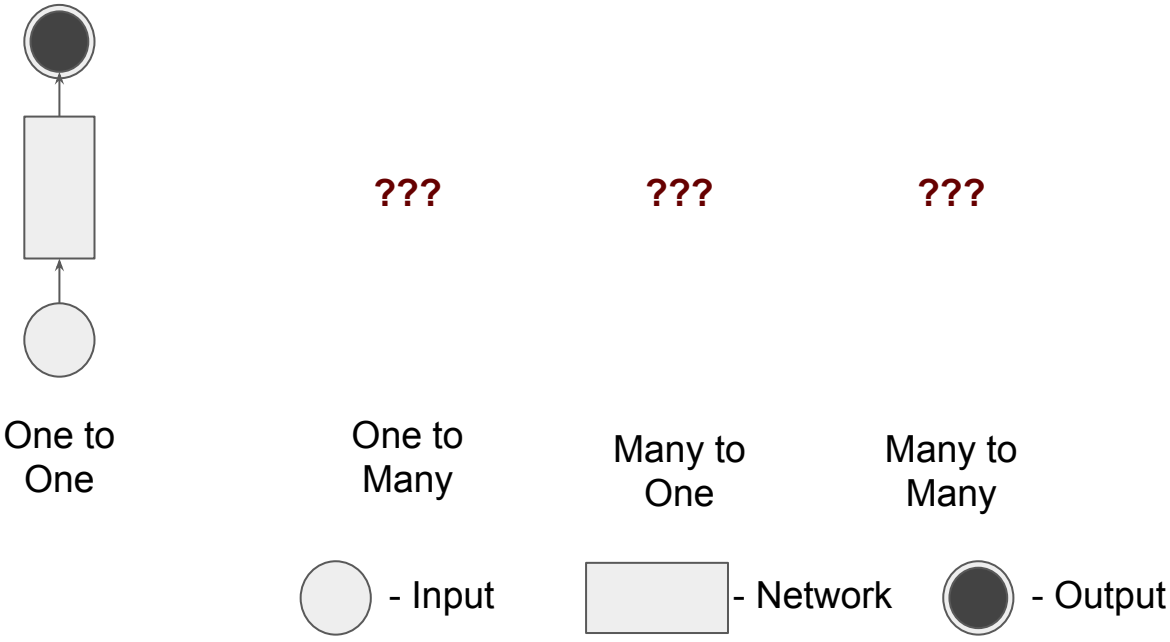

Neural Networks

Input in relation with the output

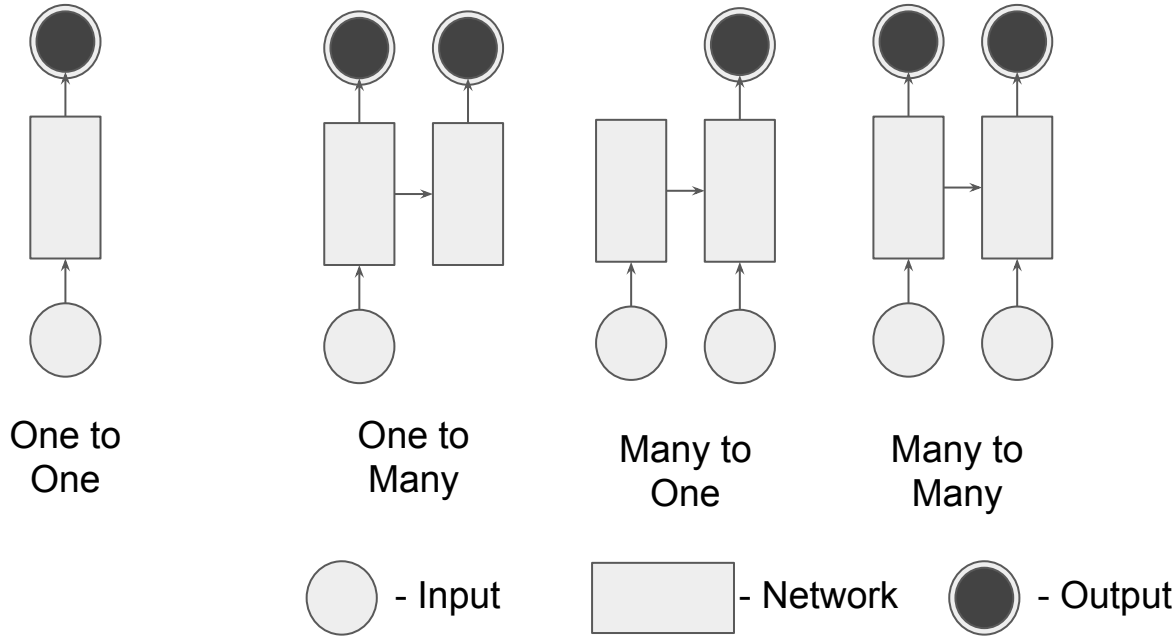

Neural Networks

Input in relation with the output

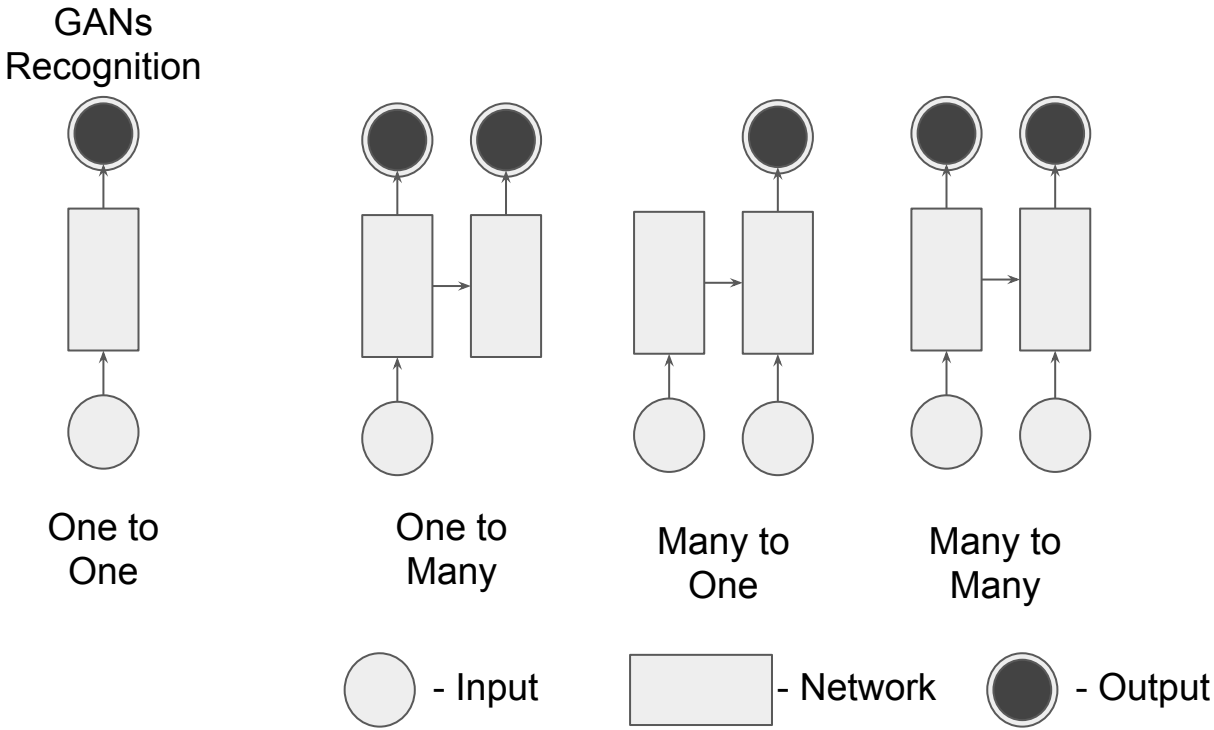

Neural Networks

Input in relation with the output

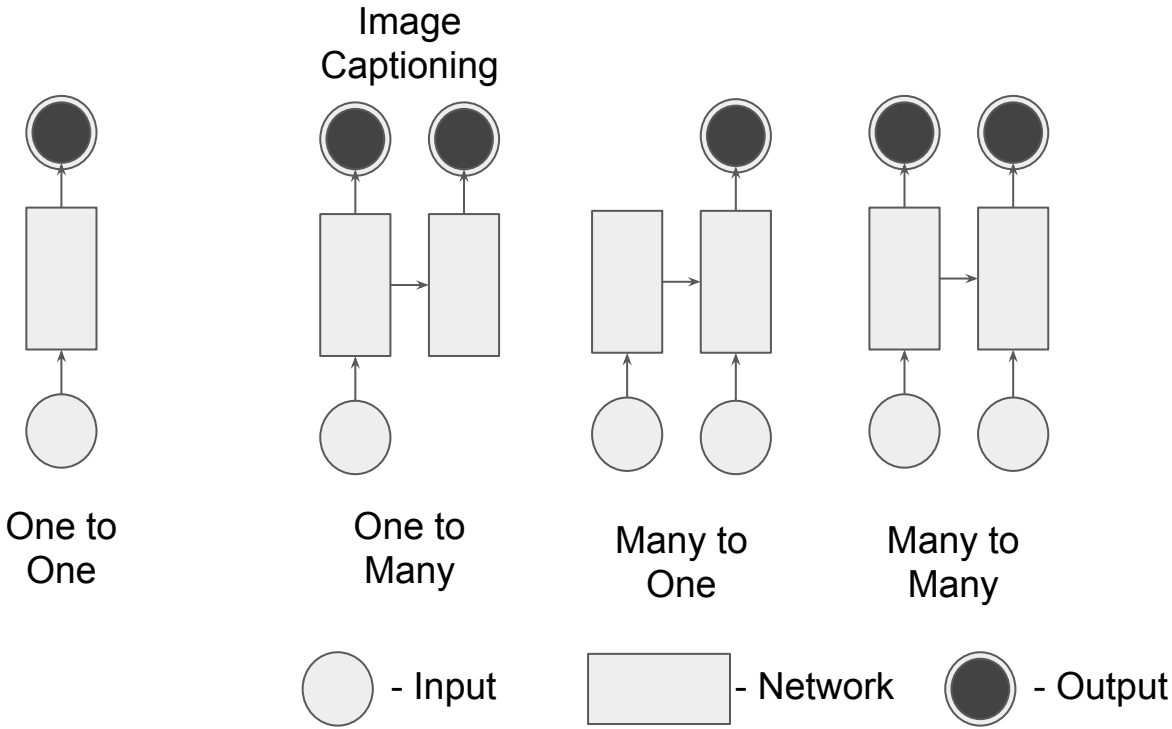

Neural Networks

Input in relation with the output

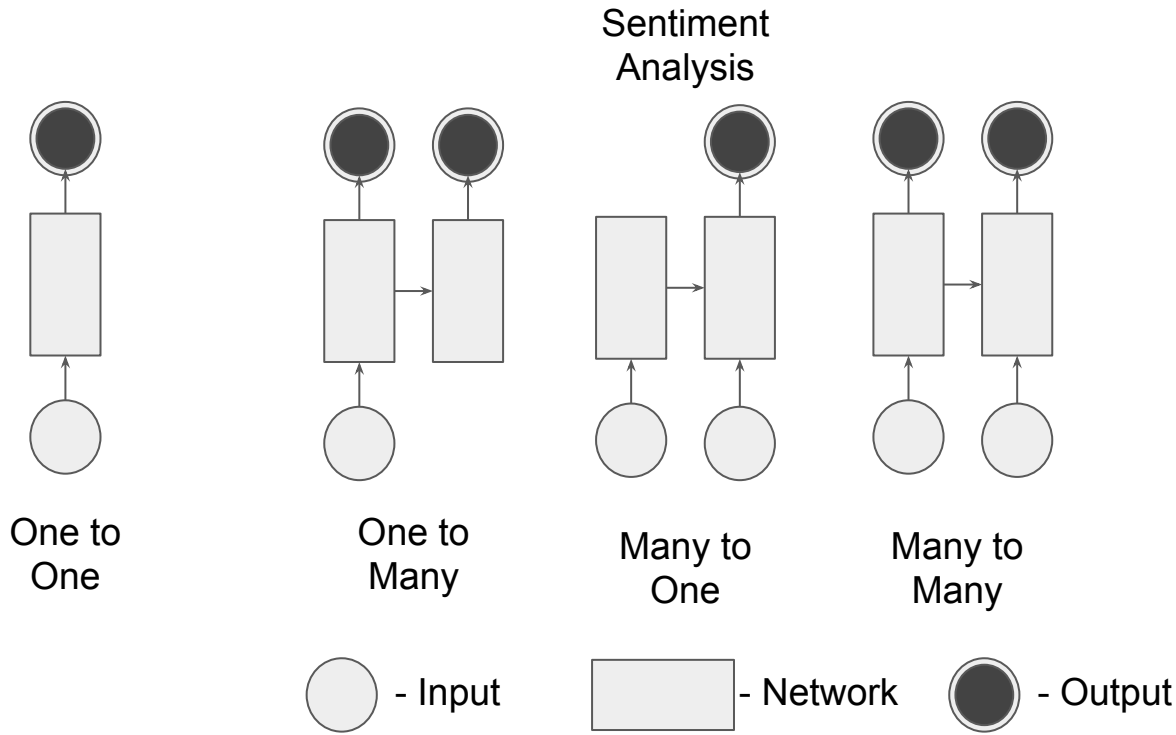

Neural Networks

Input in relation with the output

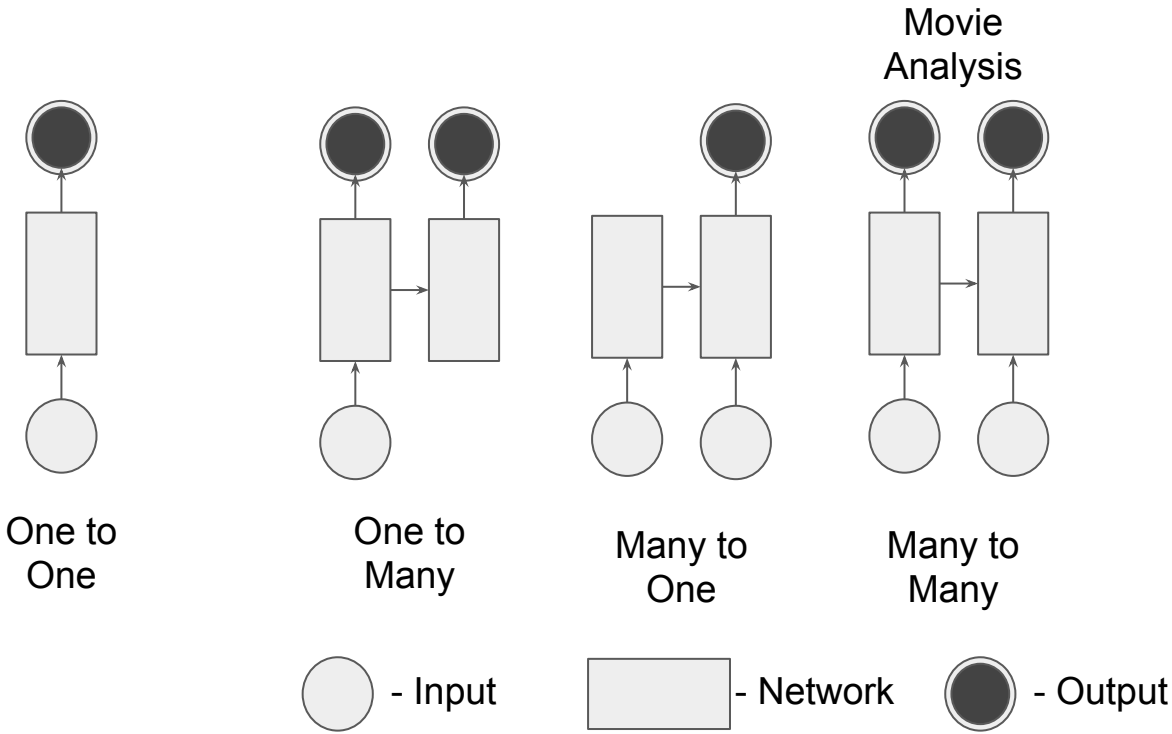

Neural Networks

Input in relation with the output

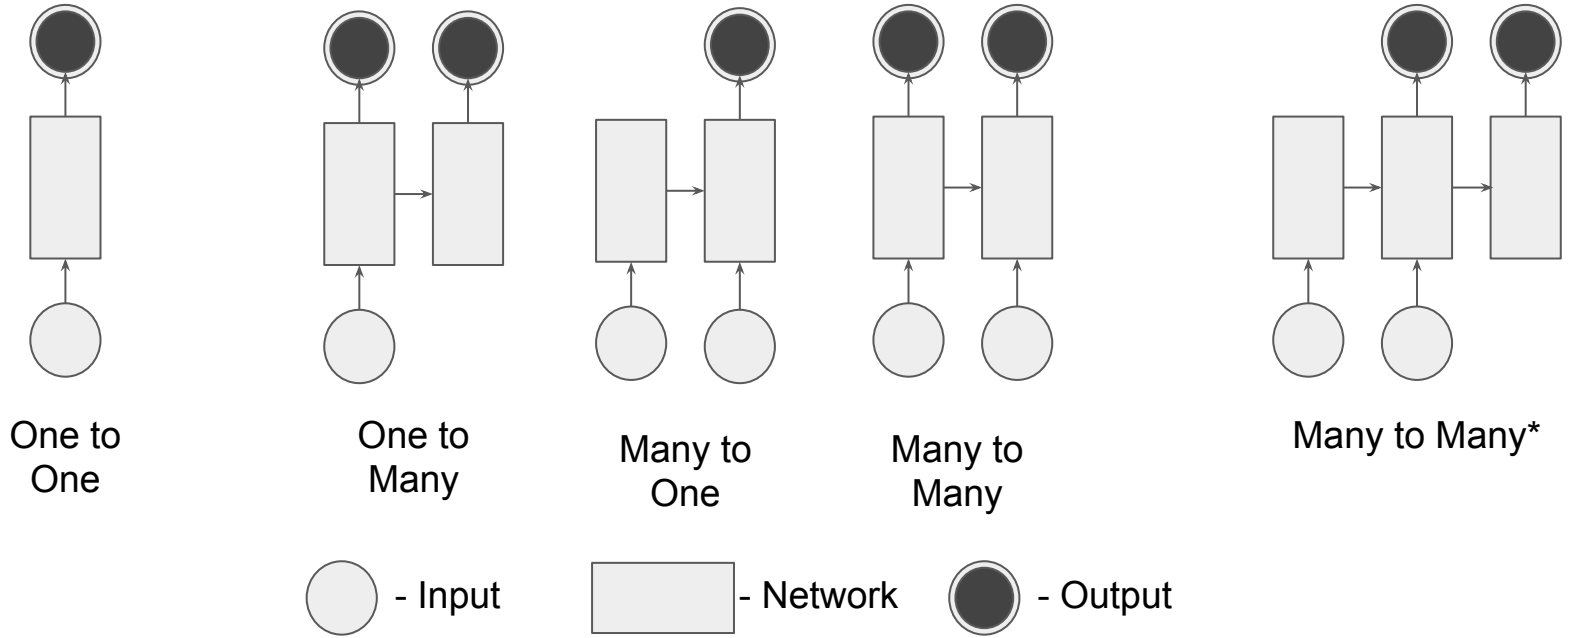

Neural Networks

Input in relation with the output

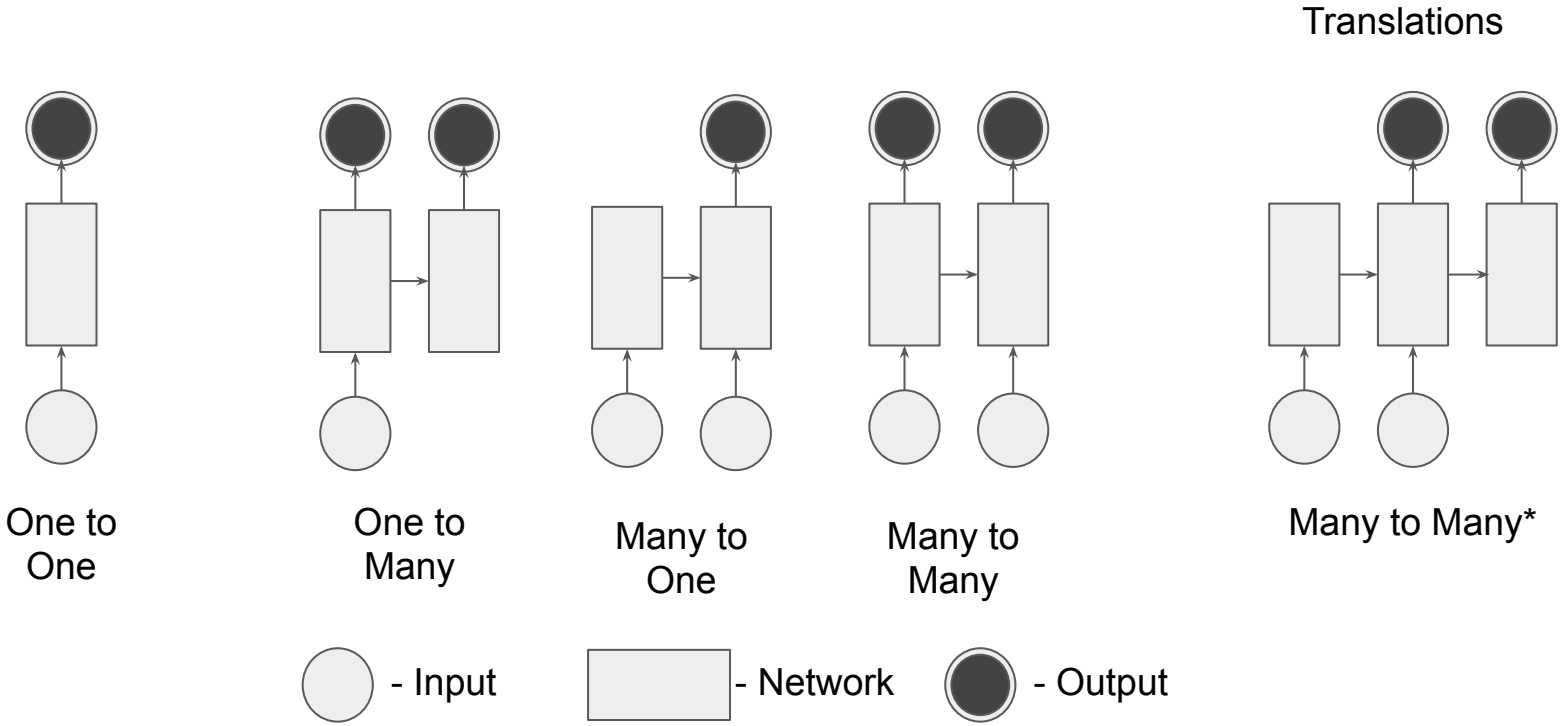

Neural Networks

Input in relation with the output

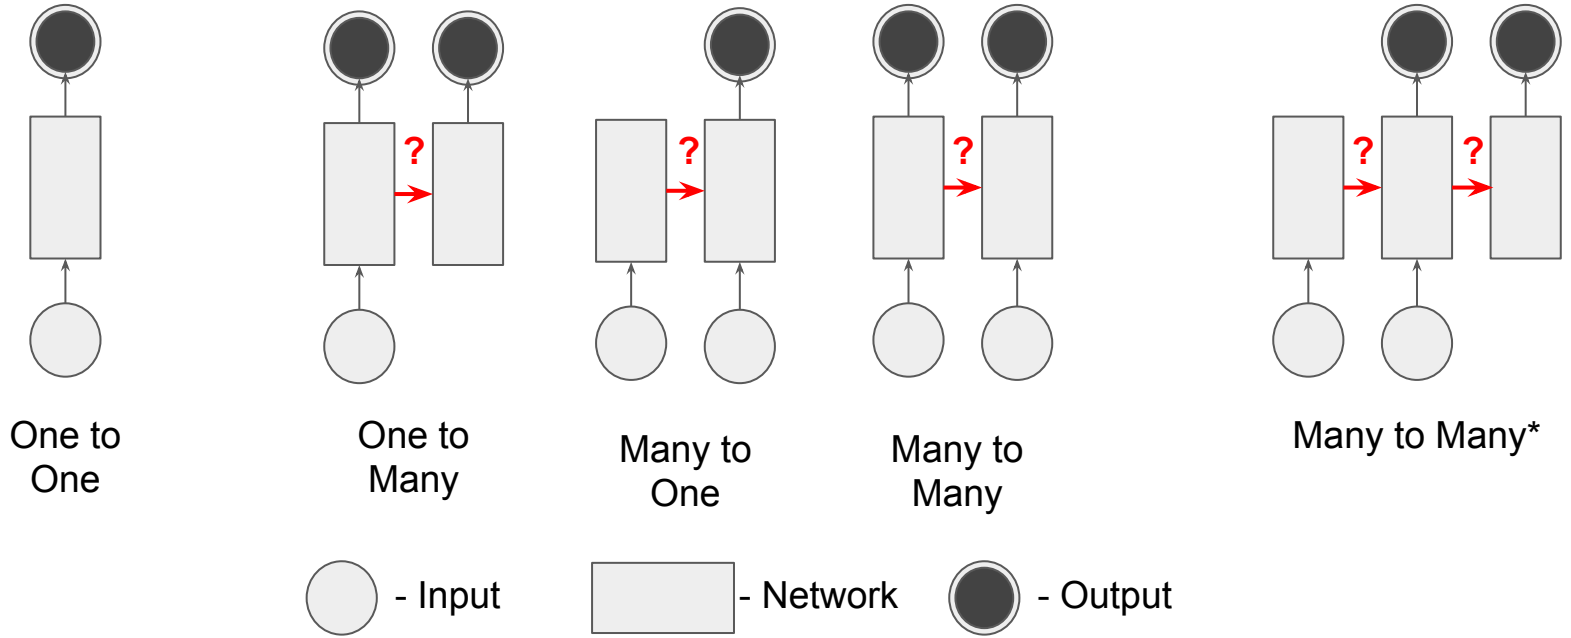

Recurrent Neural Networks

Input in relation with the output

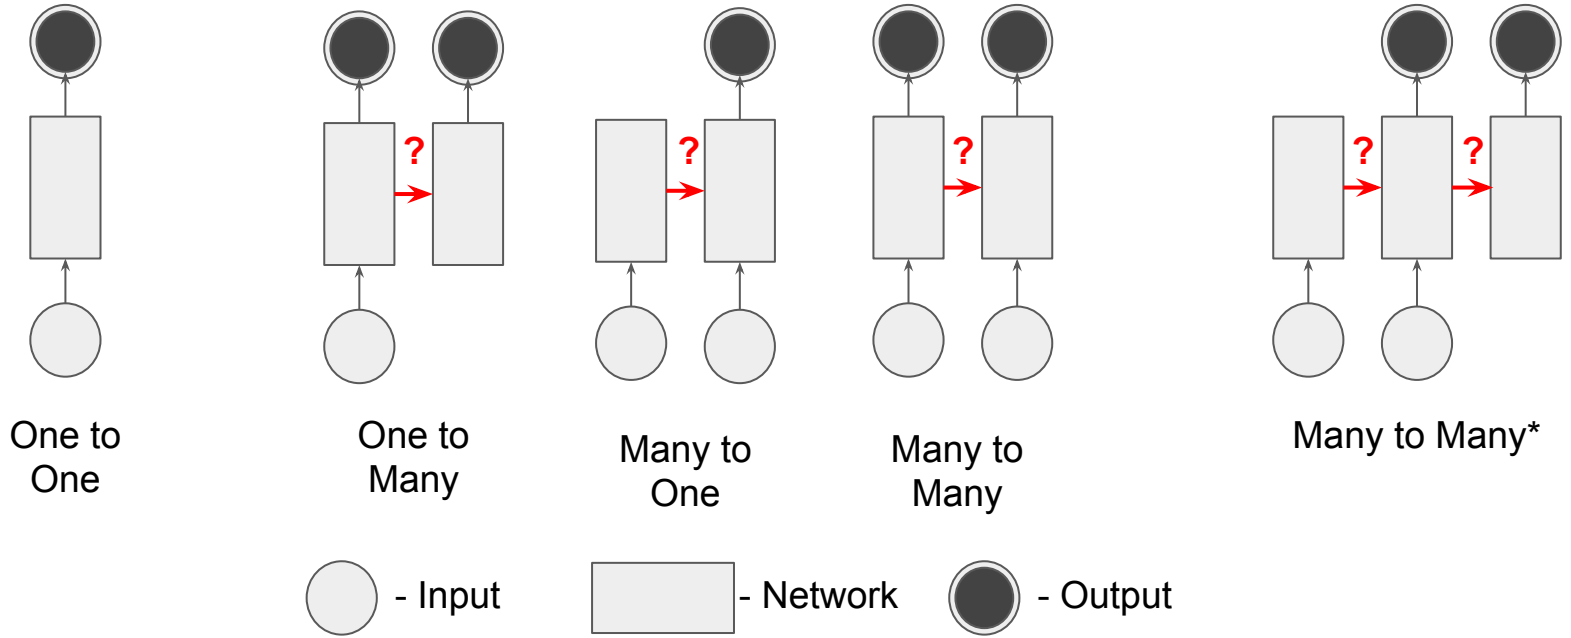

Recurrent Neural Networks

Input in relation with the output

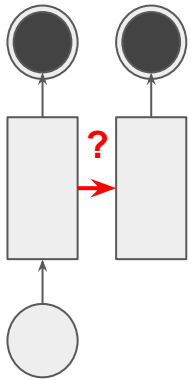

Recurrent Neural Networks

Input in relation with the output

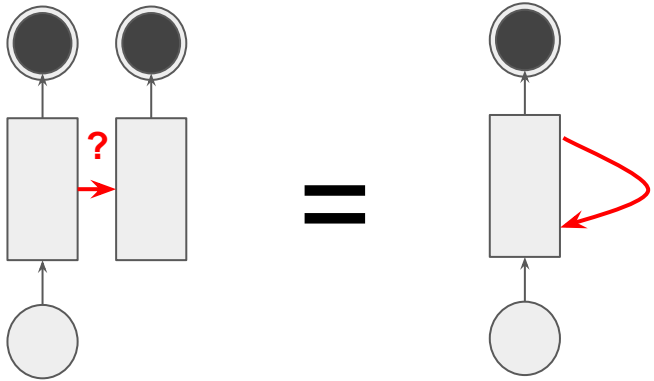

Recurrent Neural Networks

Input in relation with the output

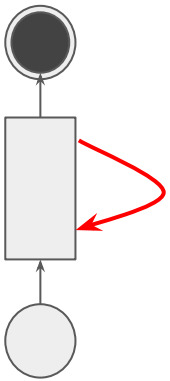

Recurrent Neural Networks

Input in relation with the output

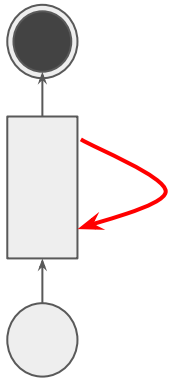

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

Explained

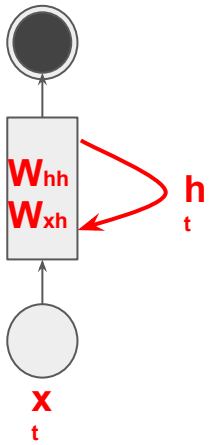

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

Explained

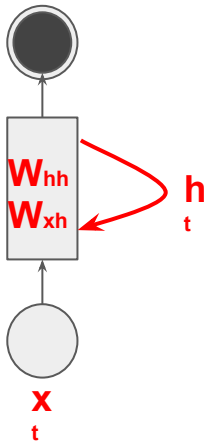

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

Explained

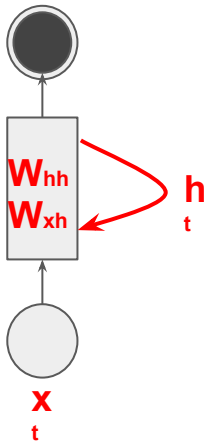

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

The problem

Recurrent Neural Networks

The problem

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

The problem

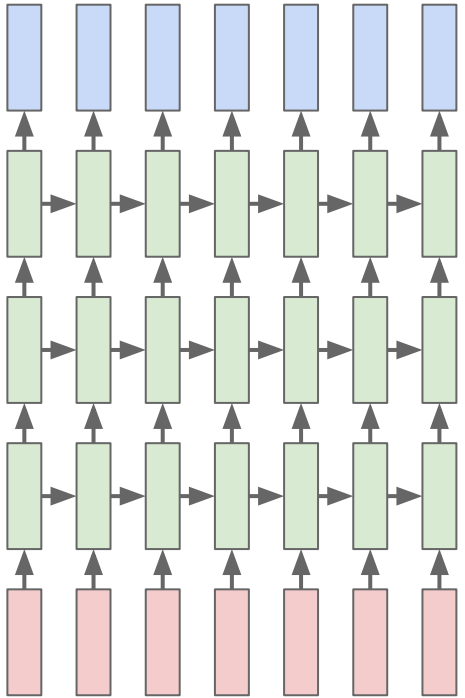

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

The problem

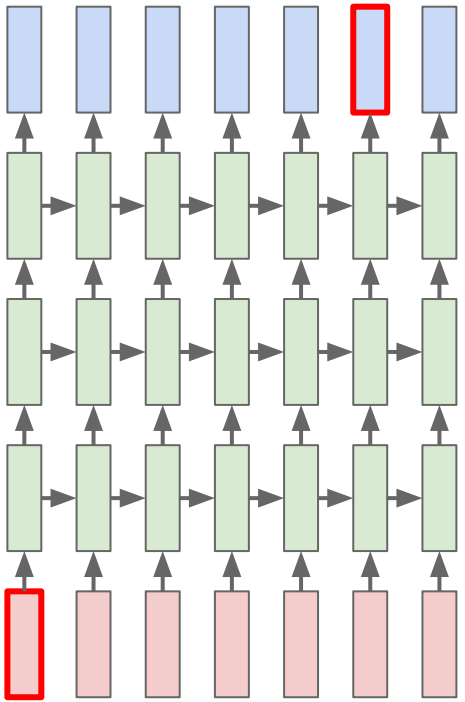

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

The problem

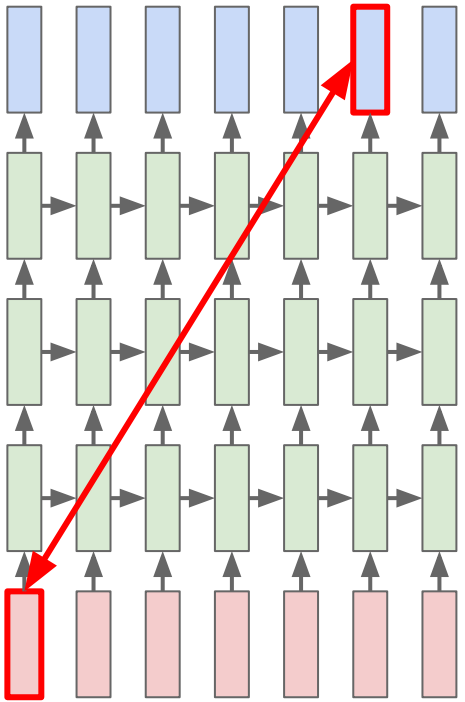

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Recurrent Neural Networks

The problem

Vanishing gradient

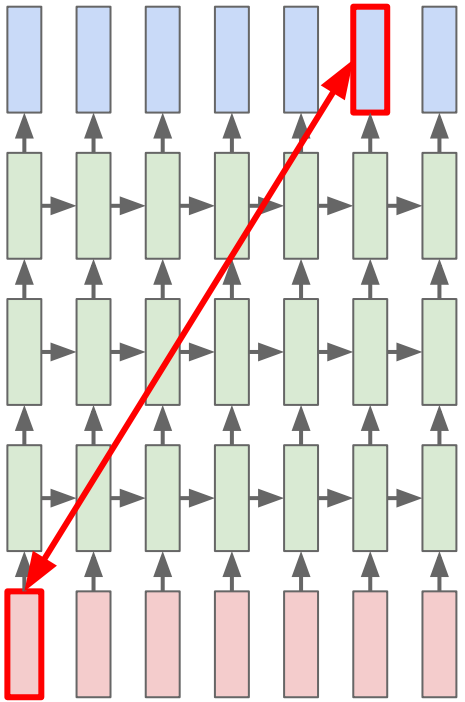

$$h_{t+1} = \tanh(W_{hh}h_t + W_{xh}x_t)$$

Long short-term memory cell

Long short-term memory cell

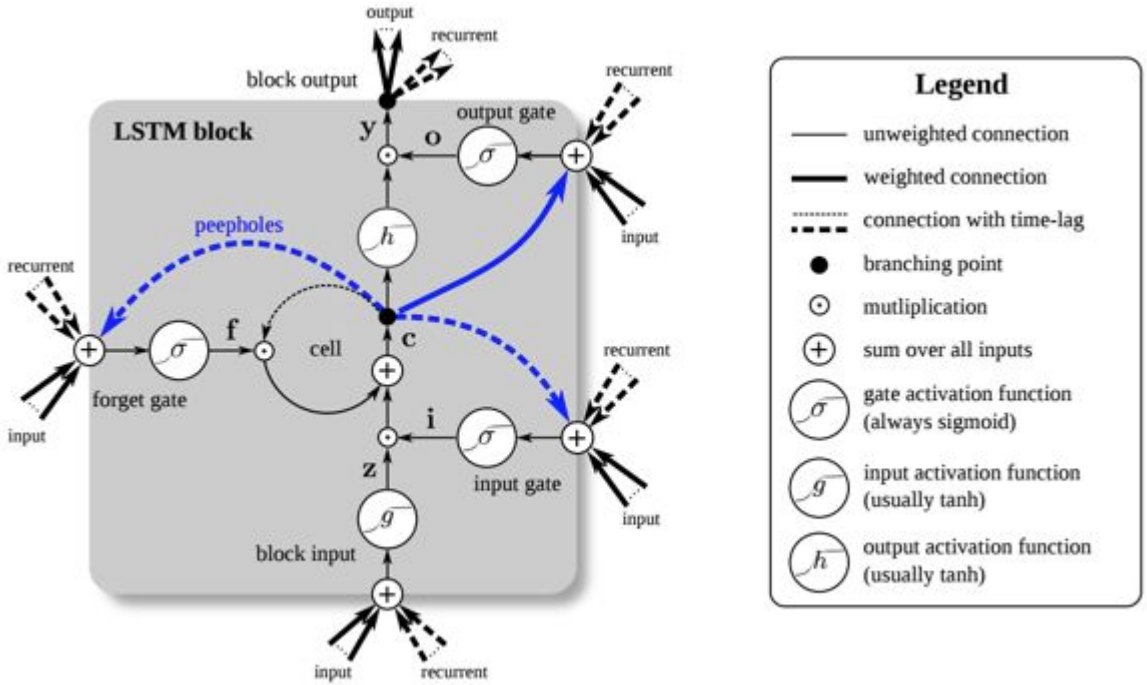

[1]

Long short-term memory cell

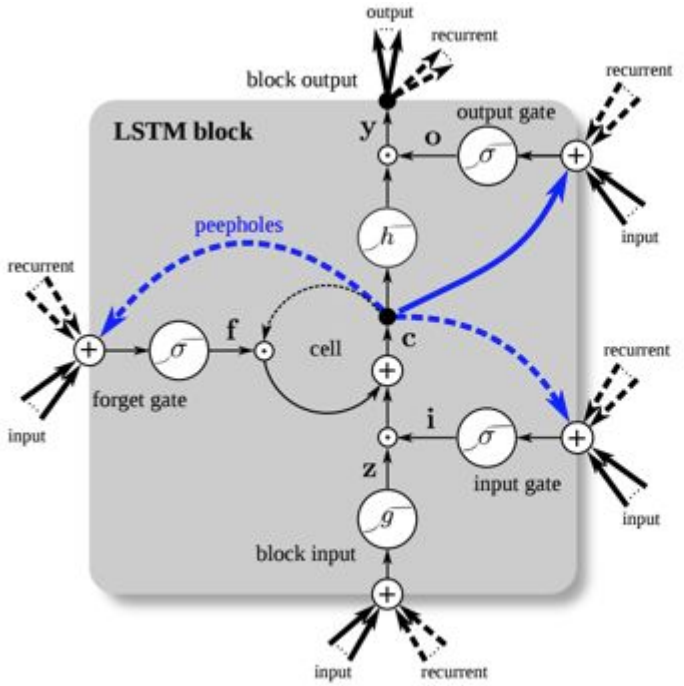

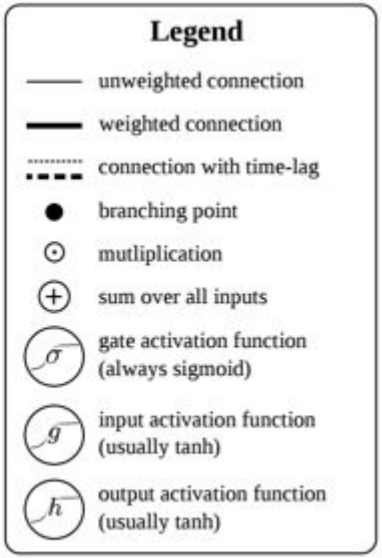

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \tanh \end{pmatrix} W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g$$

$$h_t^l = o \odot \tanh(c_t^l)$$

[1]

Code example

[Github]

Sources

[1] “LSTM: A Search Space Odyssey” -- https://arxiv.org/pdf/1503.04069.pdf

[2] “RNN Escapades” -- London ML meetup 09/2015 Andrej Karpathy
https://docs.google.com/presentation/d/1qs2IuSdZvbNfzw217kH5-1Z9DjG0Ng6fJiabaLNQVaY